

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5 import seaborn as sns
6 from IPython import get_ipython
7 import warnings
8 warnings.filterwarnings("ignore")
```

In [2]:

```
1 data = pd.read_csv('hr_data.csv')
```

In [3]:

```
1 data.head()
```

Out[3]:

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company
0	0.38	0.53	2	157	3
1	0.80	0.86	5	262	6
2	0.11	0.88	7	272	4
3	0.72	0.87	5	223	1
4	0.37	0.52	2	159	3

In [4]:

```
1 data.tail()
```

Out[4]:

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company
14994	0.40	0.57	2	151	3
14995	0.37	0.48	2	160	3
14996	0.37	0.53	2	143	3
14997	0.11	0.96	6	280	3
14998	0.37	0.52	2	158	3

In [5]:



```
1 data.shape
```

Out[5]:

```
(14999, 10)
```

In [6]:



```
1 data.columns
```

Out[6]:

```
Index(['satisfaction_level', 'last_evaluation', 'number_project',  
      'average_monthly_hours', 'time_spend_company', 'Work_accident', 'left',  
      'promotion_last_5years', 'sales', 'salary'],  
      dtype='object')
```

In [7]:



```
1 data.duplicated().sum()
```

Out[7]:

```
3008
```

In [8]:




```
1 data.drop_duplicates()
```

Out[8]:

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_con
0	0.38	0.53	2	157	
1	0.80	0.86	5	262	
2	0.11	0.88	7	272	
3	0.72	0.87	5	223	
4	0.37	0.52	2	159	
...
11995	0.90	0.55	3	259	
11996	0.74	0.95	5	266	
11997	0.85	0.54	3	185	
11998	0.33	0.65	3	172	
11999	0.50	0.73	4	180	

```
11991 rows × 10 columns
```



In [9]:



```
1 data.isnull().sum()
```

Out[9]:

```
satisfaction_level      0
last_evaluation         0
number_project          0
average_monthly_hours   0
time_spend_company      0
work_accident           0
left                   0
promotion_last_5years   0
sales                   0
salary                  0
dtype: int64
```

In [10]:



```
1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   satisfaction_level      14999 non-null  float64
1   last_evaluation        14999 non-null  float64
2   number_project         14999 non-null  int64
3   average_monthly_hours  14999 non-null  int64
4   time_spend_company     14999 non-null  int64
5   work_accident          14999 non-null  int64
6   left                   14999 non-null  int64
7   promotion_last_5years  14999 non-null  int64
8   sales                  14999 non-null  object
9   salary                  14999 non-null  object
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB
```

In [11]:



```
1 data.describe()
```

Out[11]:

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company
count	14999.000000	14999.000000	14999.000000	14999.000000	14999.000000
mean	0.612834	0.716102	3.803054	201.050337	3.462051
std	0.248631	0.171169	1.232592	49.943099	1.460209
min	0.090000	0.360000	2.000000	96.000000	2.000000
25%	0.440000	0.560000	3.000000	156.000000	3.000000
50%	0.640000	0.720000	4.000000	200.000000	3.000000
75%	0.820000	0.870000	5.000000	245.000000	4.000000
max	1.000000	1.000000	7.000000	310.000000	10.000000

In [12]:



```
1 data.nunique()
```

Out[12]:

```
satisfaction_level      92
last_evaluation          65
number_project           6
average_monthly_hours   215
time_spend_company       8
Work_accident           2
left                     2
promotion_last_5years    2
sales                    10
salary                   3
dtype: int64
```

In [13]:



```
1 data['number_project'].unique()
```

Out[13]:

```
array([2, 5, 7, 6, 4, 3], dtype=int64)
```

In [14]:

```
1 data['number_project'].value_counts()
```

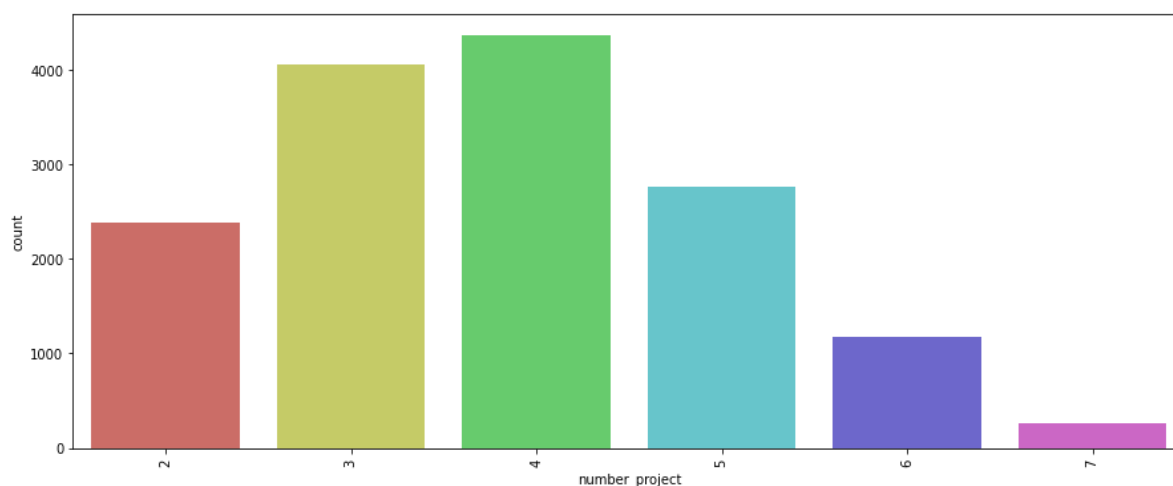
Out[14]:

```
4    4365
3    4055
5    2761
2    2388
6    1174
7     256
```

Name: number_project, dtype: int64

In [15]:

```
1 plt.figure(figsize=(15,6))
2 sns.countplot('number_project', data = data, palette = 'hls')
3 plt.xticks(rotation = 90)
4 plt.show()
```



In [16]:

```
1 data['time_spend_company'].unique()
```

Out[16]:

```
array([ 3,  6,  4,  5,  2,  8, 10,  7], dtype=int64)
```

In [17]:

```
1 data['time_spend_company'].value_counts()
```


Out[17]:

```
3    6443
2    3244
4    2557
5    1473
6     718
10    214
7     188
8     162
```

Name: time_spend_company, dtype: int64

In [18]:

```
1 plt.figure(figsize=(15,6))
2 sns.countplot('time_spend_company', data = data, palette = 'hls')
3 plt.xticks(rotation = 90)
4 plt.show()
```



In [19]:

```
1 data['Work_accident'].unique()
```

Out[19]:

```
array([0, 1], dtype=int64)
```

In [20]:

```
1 data['Work_accident'].value_counts()
```

Out[20]:

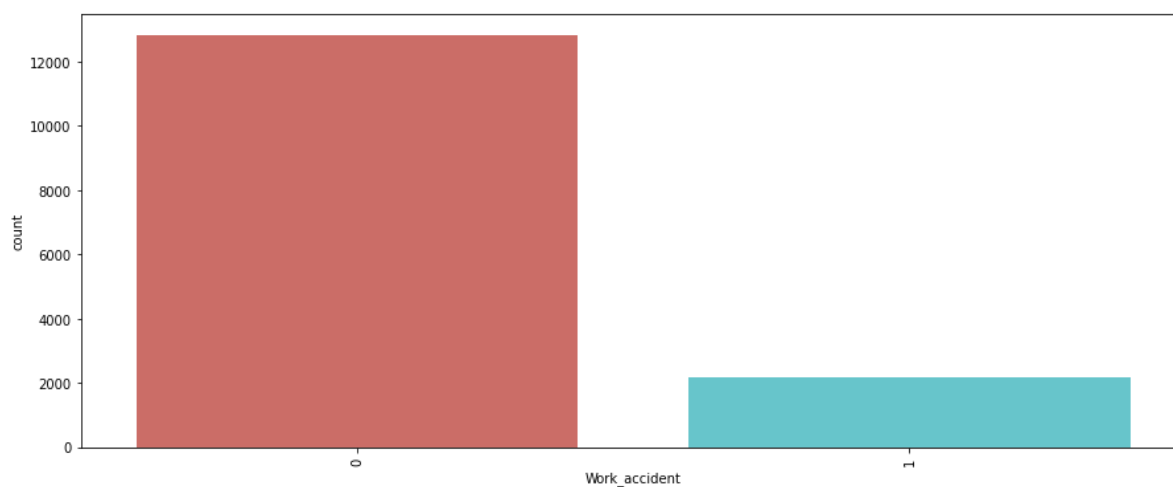
```
0    12830
1     2169
```

Name: Work_accident, dtype: int64

In [21]:



```
1 plt.figure(figsize=(15,6))
2 sns.countplot('Work_accident', data = data, palette = 'hls')
3 plt.xticks(rotation = 90)
4 plt.show()
```



In [22]:



```
1 data['left'].unique()
```

Out[22]:

```
array([1, 0], dtype=int64)
```

In [23]:



```
1 data['left'].value_counts()
```

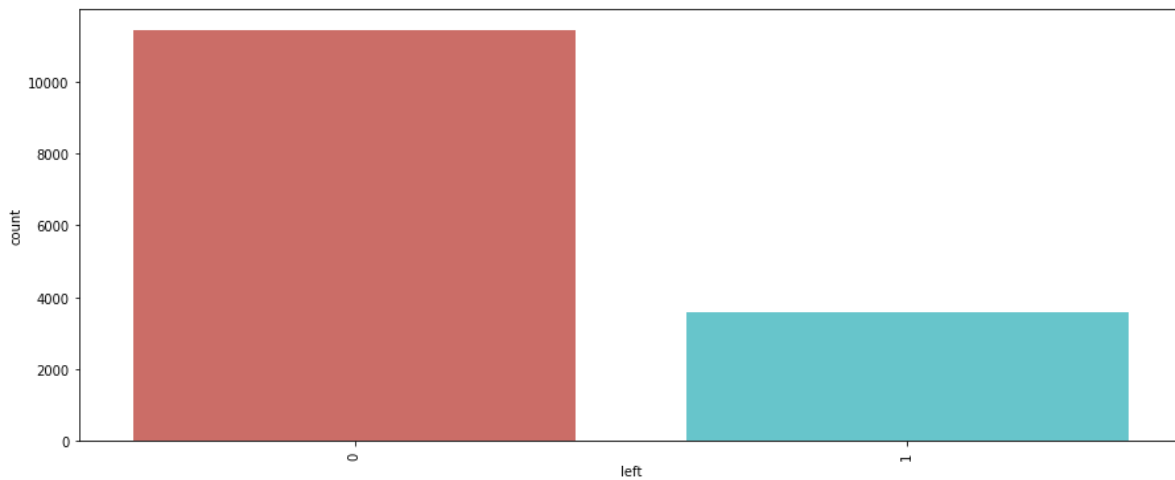
Out[23]:

```
0    11428
1     3571
Name: left, dtype: int64
```

In [24]:



```
1 plt.figure(figsize=(15,6))
2 sns.countplot('left', data = data, palette = 'hls')
3 plt.xticks(rotation = 90)
4 plt.show()
```



In [25]:



```
1 data['promotion_last_5years'].unique()
```

Out[25]:

```
array([0, 1], dtype=int64)
```

In [26]:



```
1 data['promotion_last_5years'].value_counts()
```

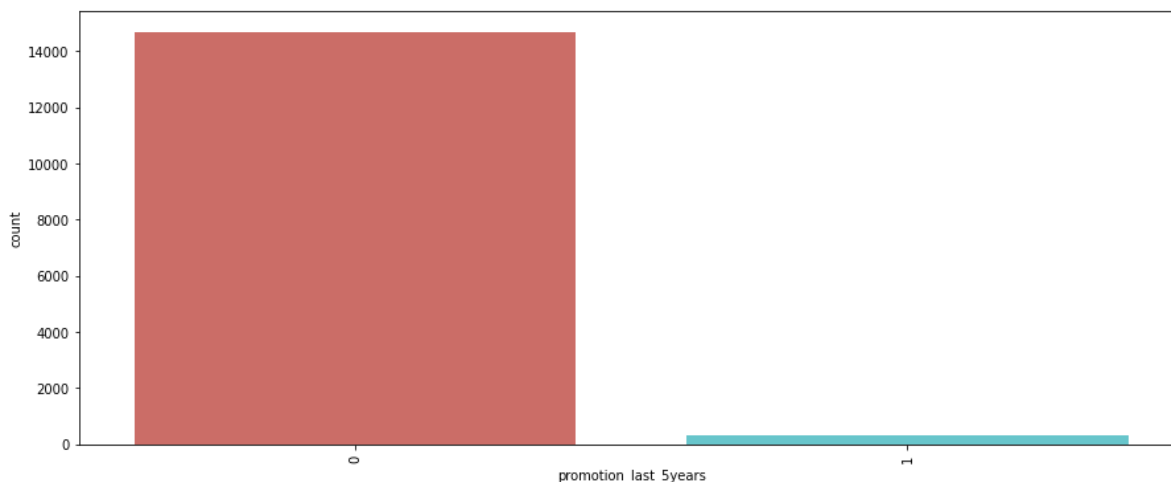
Out[26]:

```
0    14680
1      319
Name: promotion_last_5years, dtype: int64
```


In [27]:



```
1 plt.figure(figsize=(15,6))
2 sns.countplot('promotion_last_5years', data = data, palette = 'hls')
3 plt.xticks(rotation = 90)
4 plt.show()
```



In [28]:



```
1 data['sales'].unique()
```

Out[28]:

```
array(['sales', 'accounting', 'hr', 'technical', 'support', 'management',
      'IT', 'product_mng', 'marketing', 'RandD'], dtype=object)
```

In [29]:



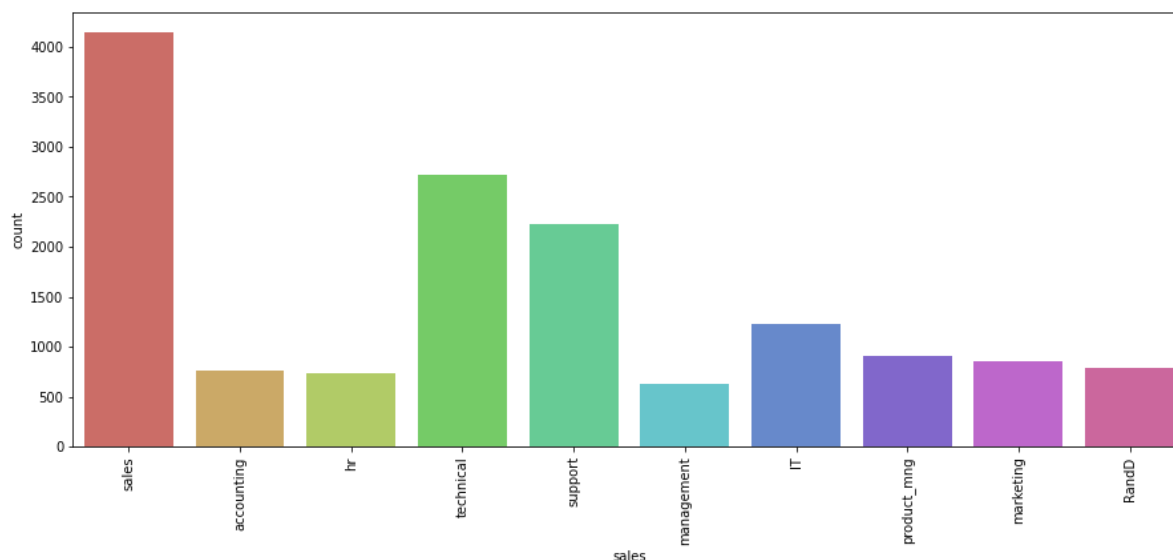
```
1 data['sales'].value_counts()
```

Out[29]:

```
sales          4140
technical      2720
support        2229
IT             1227
product_mng    902
marketing       858
RandD          787
accounting     767
hr             739
management     630
Name: sales, dtype: int64
```

In [30]:

```
1 plt.figure(figsize=(15,6))
2 sns.countplot('sales', data = data, palette = 'hls')
3 plt.xticks(rotation = 90)
4 plt.show()
```



In [31]:

```
1 data['salary'].unique()
```

Out[31]:

```
array(['low', 'medium', 'high'], dtype=object)
```

In [32]:

```
1 data['salary'].value_counts()
```


Out[32]:

```
low      7316
medium   6446
high     1237
Name: salary, dtype: int64
```

In [33]:



```
1 plt.figure(figsize=(15,6))
2 sns.countplot('salary', data = data, palette = 'hls')
3 plt.xticks(rotation = 90)
4 plt.show()
```



In [34]:



```
1 data = data.rename(columns = {'sales':'department'})
```

In [35]:



```
1 data['department']=np.where(data['department'] == 'support',
2                             'technical', data['department'])
3 data['department']=np.where(data['department'] == 'IT',
4                             'technical', data['department'])
```

In [36]:



```
1 cat_vars=['department','salary']
2 for var in cat_vars:
3     cat_list='var'+ '_' +var
4     cat_list = pd.get_dummies(data[var], prefix=var)
5     data1=data.join(cat_list)
6     data=data1
```

In [37]:

```
1 data.drop(data.columns[[8, 9]], axis=1, inplace=True)
2 data.columns.values
```

Out[37]:

```
array(['satisfaction_level', 'last_evaluation', 'number_project',
      'average_monthly_hours', 'time_spend_company', 'Work_accident',
      'left', 'promotion_last_5years', 'department_RandD',
      'department_accounting', 'department_hr', 'department_management',
      'department_marketing', 'department_product_mng',
      'department_sales', 'department_technical', 'salary_high',
      'salary_low', 'salary_medium'], dtype=object)
```

In [38]:

```
1 hr_vars=data.columns.values.tolist()
2 y=['left']
3 X=[i for i in hr_vars if i not in y]
```

In [41]:

```
1 cols=['satisfaction_level', 'last_evaluation', 'time_spend_company', 'Work_accident',
2       'department_RandD', 'department_hr', 'department_management', 'salary_high',
3       X=data[cols]
4       y=data['left']
```

In [43]:

```
1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(X, y,
3                                                    test_size=0.3,
4                                                    random_state=0)
5 from sklearn.linear_model import LogisticRegression
6 from sklearn import metrics
7 logreg = LogisticRegression()
8 logreg.fit(X_train, y_train)
```

Out[43]:

LogisticRegression()

In [44]:

```
1 from sklearn.metrics import accuracy_score
2 print('Logistic regression accuracy: {:.3f}'.format(accuracy_score(y_test,
3                                                                    logreg.predict(X
```

Logistic regression accuracy: 0.771

In [45]:



```
1 from sklearn.ensemble import RandomForestClassifier
2 rf = RandomForestClassifier()
3 rf.fit(X_train, y_train)
```

Out[45]:

RandomForestClassifier()

In [46]:



```
1 print('Random Forest Accuracy: {:.3f}'.format(accuracy_score(y_test,
2                                                                rf.predict(X_test))))
```

Random Forest Accuracy: 0.979

In [47]:



```
1 from sklearn.metrics import classification_report
2 print(classification_report(y_test, rf.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.99	0.98	0.99	3462
1	0.95	0.96	0.95	1038
accuracy			0.98	4500
macro avg	0.97	0.97	0.97	4500
weighted avg	0.98	0.98	0.98	4500

In [49]:

```

1 y_pred = rf.predict(X_test)
2 from sklearn.metrics import confusion_matrix
3 import seaborn as sns
4 forest_cm = metrics.confusion_matrix(y_pred, y_test)
5 sns.heatmap(forest_cm, annot=True, fmt='.2f', xticklabels = ["Left", "Stayed"] , yti
6 plt.ylabel('True class')
7 plt.xlabel('Predicted class')
8 plt.title('Random Forest')

```

Out[49]:

Text(0.5, 1.0, 'Random Forest')



In [50]:

```

1 print(classification_report(y_test, logreg.predict(X_test)))

```

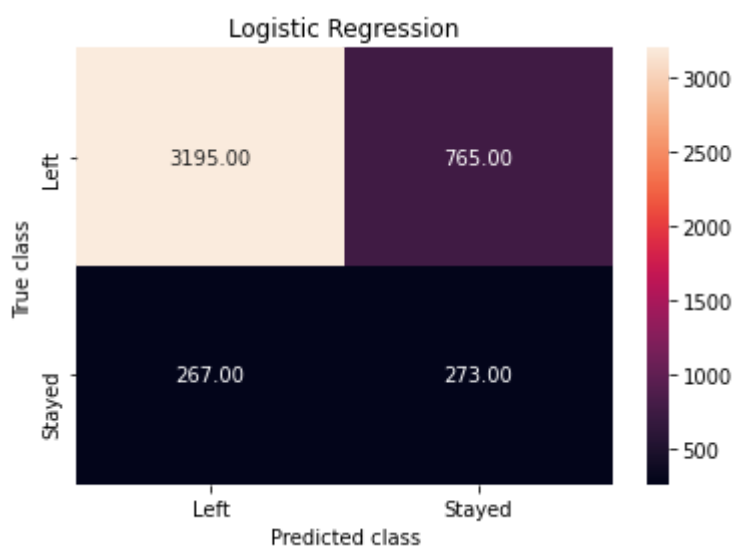
	precision	recall	f1-score	support
0	0.81	0.92	0.86	3462
1	0.51	0.26	0.35	1038
accuracy			0.77	4500
macro avg	0.66	0.59	0.60	4500
weighted avg	0.74	0.77	0.74	4500

In [52]:

```
1 logreg_y_pred = logreg.predict(X_test)
2 logreg_cm = metrics.confusion_matrix(logreg_y_pred, y_test)
3 sns.heatmap(logreg_cm, annot=True, fmt='.2f',
4             xticklabels = ["Left", "Stayed"] ,
5             yticklabels = ["Left", "Stayed"] )
6 plt.ylabel('True class')
7 plt.xlabel('Predicted class')
8 plt.title('Logistic Regression')
```

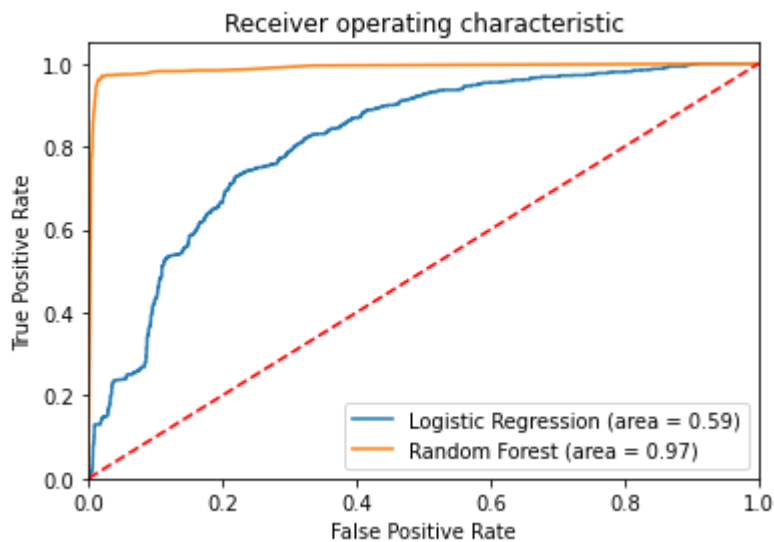
Out[52]:

Text(0.5, 1.0, 'Logistic Regression')



In [53]:

```
1 from sklearn.metrics import roc_auc_score
2 from sklearn.metrics import roc_curve
3 logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))
4 fpr, tpr, thresholds = roc_curve(y_test, logreg.predict_proba(X_test)[:,1])
5 rf_roc_auc = roc_auc_score(y_test, rf.predict(X_test))
6 rf_fpr, rf_tpr, rf_thresholds = roc_curve(y_test, rf.predict_proba(X_test)[:,1])
7 plt.figure()
8 plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
9 plt.plot(rf_fpr, rf_tpr, label='Random Forest (area = %0.2f)' % rf_roc_auc)
10 plt.plot([0, 1], [0, 1], 'r--')
11 plt.xlim([0.0, 1.0])
12 plt.ylim([0.0, 1.05])
13 plt.xlabel('False Positive Rate')
14 plt.ylabel('True Positive Rate')
15 plt.title('Receiver operating characteristic')
16 plt.legend(loc="lower right")
17 plt.show()
```



In [54]:



```
1 feature_labels = np.array(['satisfaction_level', 'last_evaluation', 'time_spend_com
2 'department_RandD', 'department_hr', 'department_management', 'salary_high',
3 importance = rf.feature_importances_
4 feature_indexes_by_importance = importance.argsort()
5 for index in feature_indexes_by_importance:
6     print('{ }-{: .2f}%'.format(feature_labels[index], (importance[index] *100.0)))
```

```
promotion_last_5years-0.22%
department_management-0.26%
department_RandD-0.29%
department_hr-0.30%
salary_high-0.63%
salary_low-1.18%
Work_accident-1.46%
last_evaluation-18.31%
time_spend_company-25.70%
satisfaction_level-51.65%
```