

```
In [55]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn import svm
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score
from sklearn.ensemble import VotingClassifier
```

```
In [3]: wine=pd.read_csv('C:\\Users\\LENOVO\\Documents\\wine.csv')
```

```
In [4]: # view the first 10 of our dataset
wine.head(10)
```

Out[4]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	9.5
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	10.5

```
In [5]: # the last 5 from our dataset
wine.tail()
```

Out[5]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11

```
In [6]: wine.shape
```

Out[6]: (1599, 12)

```
In [7]: #statistical decription of the dataset
wine.describe()
```

Out[7]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000

In [8]: `wine.isnull()`

Out[8]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcol
0	False	False	False	False	False	False	False	False	False	False	Fa
1	False	False	False	False	False	False	False	False	False	False	Fa
2	False	False	False	False	False	False	False	False	False	False	Fa
3	False	False	False	False	False	False	False	False	False	False	Fa
4	False	False	False	False	False	False	False	False	False	False	Fa
...	...	...	...	...	...	...	...	...	...	...	
1594	False	False	False	False	False	False	False	False	False	False	Fa
1595	False	False	False	False	False	False	False	False	False	False	Fa
1596	False	False	False	False	False	False	False	False	False	False	Fa
1597	False	False	False	False	False	False	False	False	False	False	Fa
1598	False	False	False	False	False	False	False	False	False	False	Fa

1599 rows × 12 columns



In [9]: `wine.columns`

Out[9]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar', 'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH', 'sulphates', 'alcohol', 'quality'], dtype='object')

In [10]: wine.nunique

```
Out[10]: <bound method DataFrame.nunique of
c acid residual sugar chlorides \
0          7.4          0.700          0.00          1.9          0.076
1          7.8          0.880          0.00          2.6          0.098
2          7.8          0.760          0.04          2.3          0.092
3         11.2          0.280          0.56          1.9          0.075
4          7.4          0.700          0.00          1.9          0.076
...          ...          ...          ...          ...          ...
1594         6.2          0.600          0.08          2.0          0.090
1595         5.9          0.550          0.10          2.2          0.062
1596         6.3          0.510          0.13          2.3          0.076
1597         5.9          0.645          0.12          2.0          0.075
1598         6.0          0.310          0.47          3.6          0.067

      free sulfur dioxide total sulfur dioxide density    pH sulphates \
0          11.0          34.0  0.99780  3.51          0.56
1          25.0          67.0  0.99680  3.20          0.68
2          15.0          54.0  0.99700  3.26          0.65
3          17.0          60.0  0.99800  3.16          0.58
4          11.0          34.0  0.99780  3.51          0.56
...          ...          ...          ...          ...          ...
1594         32.0          44.0  0.99490  3.45          0.58
1595         39.0          51.0  0.99512  3.52          0.76
1596         29.0          40.0  0.99574  3.42          0.75
1597         32.0          44.0  0.99547  3.57          0.71
1598         18.0          42.0  0.99549  3.39          0.66


      alcohol  quality
0          9.4        5
1          9.8        5
2          9.8        5
3          9.8        6
4          9.4        5
...          ...        ...
1594        10.5        5
1595        11.2        6
1596        11.0        6
1597        10.2        5
1598        11.0        6
```

[1599 rows x 12 columns]>

In [11]: wine['quality'].value\_counts()

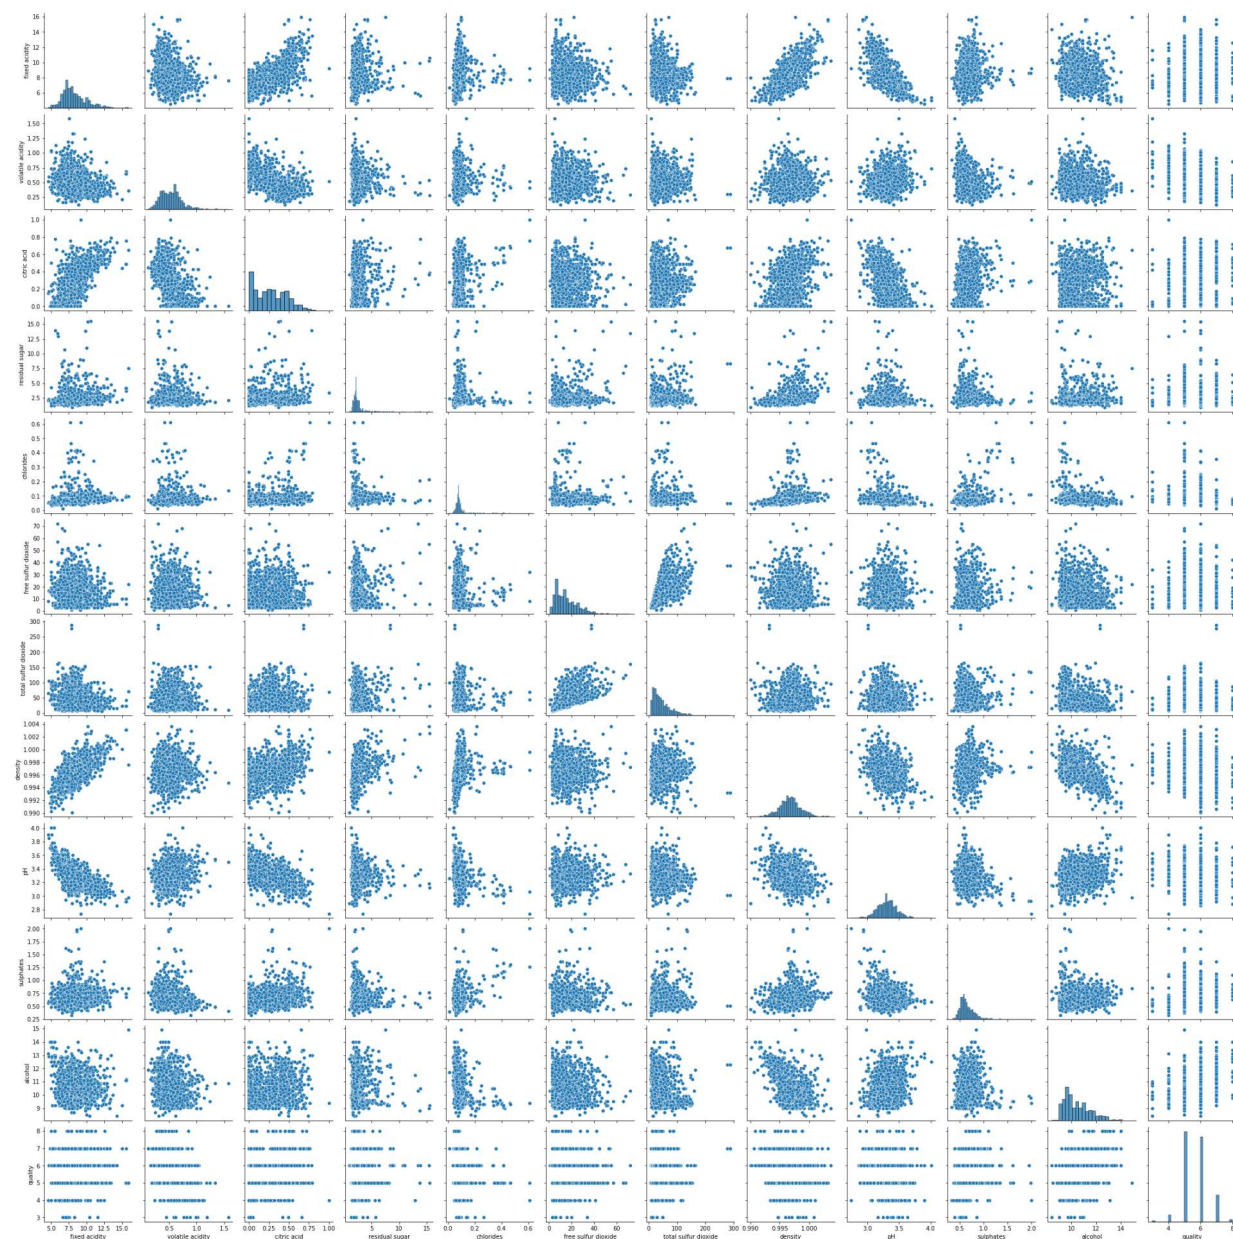
```
Out[11]: 5    681
6    638
7    199
4     53
8     18
3     10
Name: quality, dtype: int64
```

```
In [12]: plt.figure(figsize=(10,6))  
sns.scatterplot(x=wine['fixed acidity'],y=wine['volatile acidity']);
```



```
In [13]: plt.figure(figsize=(10,6))
sns.pairplot(wine);
```

<Figure size 720x432 with 0 Axes>



```
In [14]: # Preprocessing the dataset
bins=(1, 6.5, 8)
group_name=['good','bad']
wine['quality']= pd.cut(wine['quality'],bins=bins,labels=group_name)
```

```
In [15]: wine['quality'].unique()
```

```
Out[15]: ['good', 'bad']
Categories (2, object): ['good' < 'bad']
```


```
In [19]: wine[:10]
```

```
Out[19]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	9.5
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	10.5

```
In [16]: plt.figure(figsize=(10,6))  
sns.countplot(wine['quality']);
```

C:\Users\LENOVO\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(



```
In [26]: label_quality=LabelEncoder()
```

```
In [28]: wine['quality'] = label_quality.fit_transform(wine['quality'])
```



In [29]: `wine[:10]`

Out[29]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	9.5
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	10.5



In [31]: `wine['quality'].value_counts()`

Out[31]:

```
1    1382
0     217
Name: quality, dtype: int64
```

In [32]: `x=wine.drop(['quality'],axis=1)`

In [33]: x

Out[33]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcoh
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9
...	...	...	...	...	...	...	...	...	...	...	...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11

1599 rows × 11 columns

In [34]: y=wine['quality']

In [36]: y

Out[36]:

0	1
1	1
2	1
3	1
4	1
...	..
1594	1
1595	1
1596	1
1597	1
1598	1

Name: quality, Length: 1599, dtype: int32

In [37]: *# training , testing and splitting the dataset*  
x\_train,x\_test,y\_train,y\_test=train\_test\_split(x,y,test\_size=0.25,random\_state=45)

```
In [40]: # applying standard scaler to get optimizer result
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test= sc.transform(x_test)
```

```
In [41]: x_train
```

```
Out[41]: array([[ -1.09852129,  1.50992721, -1.38824517, ...,  1.14517889,
        -0.8026406 ,  1.10997762],
       [ -1.09852129,  0.6327804 , -0.30279894, ...,  1.79318941,
         0.02216045, -0.59560931],
       [ 1.33427414,  1.70799262,  0.52420772, ..., -0.99325583,
        -0.8026406 , -0.50085448],
       ...,
       [-0.75097909,  0.18005946, -0.71630226, ...,  0.17316311,
        -0.39024008, -1.16413829],
       [ 0.29164752, -0.046301 ,  0.36914398, ..., -0.34524531,
        -0.27241136, -0.8798738 ],
       [ 1.21842673, -1.51764404,  1.0927748 , ...,  0.10836206,
         1.25936202,  0.44669381]])
```

```
In [42]: #USING DECISION TREE CLASSIFIER ALGORITHM
dtc=DecisionTreeClassifier(random_state=45)
```

```
In [43]: dtc.fit(x_train,y_train)
```

```
Out[43]: DecisionTreeClassifier(random_state=45)
```

```
In [44]: dtc_pred= dtc.predict(x_test)
```

```
In [48]: dtc_pred[:30]
```

```
Out[48]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 0, 1, 1, 0, 1, 1])
```

```
In [58]: # to determine the accuracy of the model
accuracy_score(y_test,dtc_pred)*100
```

```
Out[58]: 86.0
```

```
In [50]: #USING RANDOM FOREST CLASSIFIER ALGORITHM
rfc=RandomForestClassifier(n_estimators=200)
```

```
In [51]: rfc.fit(x_train,y_train)
```

```
Out[51]: RandomForestClassifier(n_estimators=200)
```

```
In [52]: rfc_pred = rfc.predict(x_test)
```

```
In [53]: rfc_pred[:30]
```

```
Out[53]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
                1, 1, 1, 1, 1, 0, 1, 1])
```

```
In [57]: # to determine the accuracy of the model
accuracy_score(y_test,rfc_pred)*100
```

```
Out[57]: 89.5
```

```
In [59]: #USING KNEAREST NIEGHBOR ALGORITHM
knn=KNeighborsClassifier(n_neighbors=5)
```

```
In [60]: knn.fit(x_train,y_train)
```

```
Out[60]: KNeighborsClassifier()
```

```
In [61]: knn_pred= knn.predict(x_test)
```

```
In [63]: knn_pred[:30]
```

```
Out[63]: array([1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1,
                1, 1, 1, 1, 1, 0, 1, 1])
```

```
In [64]: # to determine the accuracy of the model
accuracy_score(y_test,rfc_pred)*100
```

```
Out[64]: 89.5
```

```
In [66]: nb=GaussianNB()
```

```
In [67]: nb.fit(x_train,y_train)
```

```
Out[67]: GaussianNB()
```

```
In [68]: nb_pred=nb.predict(x_test)
```

```
In [70]: nb_pred[:30]
```

```
Out[70]: array([1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1,
                1, 1, 1, 1, 1, 0, 1, 1])
```

```
In [71]: # to determine the accuracy of the model
accuracy_score(y_test,nb_pred)*100
```

```
Out[71]: 82.0
```

```
In [74]: # USING VOTING ENSEMBLE TO DECIDE THE BEST MODEL
estimators=[('Decision',dtc), ('Random',rfc),('KNeighbors',knn),('Gaussian',nb)]
```

```
In [75]: VC= VotingClassifier(estimators=estimators,voting='hard')
```

```
In [76]: VC.fit(x_train,y_train)
```

```
Out[76]: VotingClassifier(estimators=[('Decision',
                                     DecisionTreeClassifier(random_state=45)),
                                     ('Random',
                                     RandomForestClassifier(n_estimators=200)),
                                     ('KNeighbors', KNeighborsClassifier()),
                                     ('Gaussian', GaussianNB())])
```

```
In [77]: vc_pred=VC.predict(x_test)
```

```
In [78]: vc_pred[:30]
```

```
Out[78]: array([1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1,
                1, 1, 1, 1, 1, 0, 1, 1])
```

```
In [79]: # to determine the accuracy of the model
accuracy_score(y_test,vc_pred)*100
```

```
Out[79]: 88.0
```

```
In [80]: #USING BAGGING ENSEMBLE
from sklearn.ensemble import BaggingClassifier
```

```
In [81]: BC= BaggingClassifier(base_estimator=dtc,n_estimators=10)
```

```
In [82]: BC.fit(x_train,y_train)
```

```
Out[82]: BaggingClassifier(base_estimator=DecisionTreeClassifier(random_state=45))
```

```
In [83]: bc_pred=BC.predict(x_test)
```

```
In [84]: bc_pred[:30]
```

```
Out[84]: array([0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
                1, 1, 1, 1, 1, 0, 1, 1])
```

```
In [85]: # to determine the accuracy of the model
accuracy_score(y_test,bc_pred)*100
```

```
Out[85]: 87.5
```

```
In [86]: # USING STACKING ENSEMBLE
from sklearn.ensemble import StackingClassifier
```

```
In [87]: SR= StackingClassifier(estimators=estimators,final_estimator=dtc)
```

```
In [88]: SR.fit(x_train,y_train)
```

```
Out[88]: StackingClassifier(estimators=[('Decision',  
                                         DecisionTreeClassifier(random_state=45)),  
                                         ('Random',  
                                          RandomForestClassifier(n_estimators=200)),  
                                         ('KNeighbors', KNeighborsClassifier()),  
                                         ('Gaussian', GaussianNB())],  
                             final_estimator=DecisionTreeClassifier(random_state=45))
```

```
In [89]: sr_pred=SR.predict(x_test)
```

```
In [90]: sr_pred[:30]
```

```
Out[90]: array([1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1,  
                1, 1, 1, 1, 1, 0, 1, 1])
```

```
In [91]: # to determine the accuracy of the model  
accuracy_score(y_test,sr_pred)*100
```

```
Out[91]: 87.5
```

```
In [92]: # USING VOTING ENSEMBLE TO DECIDE THE BEST MODEL  
from sklearn.ensemble import VotingClassifier
```

```
In [93]: estimators=[('Voting',VC), ('Bagging',BC),('Stacking',SR)]
```

```
In [94]: VC= VotingClassifier(estimators=estimators,voting='hard')
```

```
In [95]: VC.fit(x_train,y_train)
```

```
Out[95]: VotingClassifier(estimators=[('Voting',
                                      VotingClassifier(estimators=[('Decision',
                                                                    DecisionTreeClassif
                                                                    ier(random_state=45)),
                                                                    ('Random',
                                                                    RandomForestClassif
                                                                    ier(n_estimators=200)),
                                                                    ('KNeighbors',
                                                                    KNeighborsClassifie
                                                                    r()),
                                                                    ('Gaussian',
                                                                    GaussianNB()))]),
                                      ('Bagging',
                                      BaggingClassifier(base_estimator=DecisionTreeClas
                                      sifier(random_state=45))),
                                      ('Stacking',
                                      StackingClassifier(estimators=[('Decision',
                                                                    DecisionTreeClass
                                                                    ('Random',
                                                                    RandomForestClass
                                                                    ier(n_estimators=200)),
                                                                    ('KNeighbors',
                                                                    KNeighborsClassif
                                                                    ier()),
                                                                    ('Gaussian',
                                                                    GaussianNB()))],
                                      final_estimator=DecisionTreeCl
                                      assifier(random_state=45)))]])
```

```
In [96]: final_pred=VC.predict(x_test)
```

```
In [97]: final_pred[:30]
```

```
Out[97]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
                1, 1, 1, 1, 1, 0, 1, 1])
```

```
In [98]: # to determine the accuracy of the model
accuracy_score(y_test,final_pred)*100
```

```
Out[98]: 88.75
```

```
In [100]: # pickle to save model
import pickle
```

```
In [101]: with open('wine.plk','wb') as f:
            pickle.dump(wine, f, protocol=pickle.HIGHEST_PROTOCOL)
```

```
In [ ]:
```

